

CS:4980

Foundations of Embedded Systems

The Asynchronous Model

Part I

Copyright 2014-26, Rajeev Alur and Cesare Tinelli.

Created by Cesare Tinelli at the University of Iowa from notes originally developed by Rajeev Alur at the University of Pennsylvania. These notes are copyrighted materials and may not be used in other course settings outside of the University of Iowa in their current form or modified form without the express written permission of one of the copyright holders. During this course, students are prohibited from selling notes to or being paid for taking notes by any person or commercial firm without the express written permission of one of the copyright holders.

Asynchronous Model

Recall: In the Synchronous Model, all components execute **in lock-step** in a sequence of (logical) rounds


In the **Asynchronous Model** instead the **speeds** at which different components execute are **independent**, or unknown

Examples:

- Processes in a distributed system
- Threads in a typical operating system

Key design challenge: how to achieve coordination

Example: Asynchronous Buffer



Input channel: `in` of type `bool`

Output channel: `out` of type `bool`

State variable: `x`; can be empty (`null`) or hold `0/1` value


Initialization: `x := null`

Input task: A_i processing input: `x := in`

Output task: A_o producing outputs:

Guard: `x != null` **Update:** `out := x ; x := null`


Example: Asynchronous Buffer



Execution Model: only **one task per step** is executed


- processing of inputs (by input tasks) is decoupled from production of outputs (by output tasks)
- A task can be executed if it is **enabled**, i.e., its guard holds or is missing
- If multiple tasks are enabled, one of them is executed non-deterministically

Possible execution:



Example: Asynchronous Increments

<code>nat x := 0 ; y := 0</code>
<code>A_x : x := x + 1</code>
<code>A_y : y := y + 1</code>



An *internal task* does not involve input or output channels

- can have guard condition and update code
- is executed in an internal action

In each step, execute, either task A_x or task A_y


Possible execution:

$(0,0) \rightarrow (1,0) \rightarrow (1,1) \rightarrow (1,2) \rightarrow (1,3) \rightarrow \dots \rightarrow (1,105) \rightarrow (2, 105) \rightarrow \dots$

For every m, n , state $\{ x := m, y := n \}$ is reachable


- *Interleaving* model of concurrency

Asynchronous Merge



Sequence of messages on output channel is an **arbitrary** merge of sequences of values on the two input channels

Asynchronous Merge




At every step, **exactly one** of the four tasks executes, provided its guard holds

Possible execution:

$([], []) \rightarrow ([5], []) \rightarrow ([5], [0]) \rightarrow ([5], []) \rightarrow ([5,6], []) \rightarrow ([5,6], [3]) \rightarrow ([6], [3]) \rightarrow \dots$
out: // // 0 // // 5

Exercise: What does this component do?



Components are now called *processes*

Asynchronous Process P

Set I of (typed) *input channels*

- Defines the set of inputs of the form $x?v$, where x is an input channel and v a value

Set O of (typed) *output channels*

- Defines the set of outputs of the form $y!v$, where y is an output channel and v a value

Set S of (typed) *state variables*

- Defines the set of states Q_s

An *initialization* $Init$

- Defines the set of initial states

Asynchronous Process P (cont.)

Set of *input tasks*, each associated with an input channel x

- Guard condition over *state variables* S
- Update code from *read-set* $S \cup \{x\}$ to *write-set* S
- Defines a set of *input actions* of the form $s - x?v \rightarrow t$

Set of *output tasks*, each associated with an output channel y

- Guard condition over *state variables* S
- Update code from *read-set* S to *write-set* $S \cup \{y\}$
- Defines a set of *output actions* of the form $s - y!v \rightarrow t$

Set of *internal tasks*

- Guard condition over *state variables* S
- Update code from *read-set* S to *write-set* S
- Defines a set of *internal actions* of the form $s - \varepsilon \rightarrow t$

Asynchronous Gates




Why design **asynchronous circuits**?

- Input can be changed even before the effect propagates through the entire circuit
- Can be faster than synchronous circuits, but design is more complex

Example: modeling a **NOT** gate

- When input changes, gate enters *unstable* state until it gets a chance to update its output value
- Later input changes in unstable state lead to a *hazard* state with unpredictable behavior

Asynchronous NOT Gate as an ESM




Possible execution:

$(stable, 0) \xrightarrow{out!0} (stable, 0) \xrightarrow{in?0} (unstable, 0) \xrightarrow{\varepsilon} (stable, 1) \xrightarrow{out!1} (stable, 1)$
 $\xrightarrow{in?1} (unstable, 1) \xrightarrow{out!1} (unstable, 1) \xrightarrow{in?0} (hazard, 1) \xrightarrow{out!0} (hazard, 1)$
 $\xrightarrow{out!1} (hazard, 1) \dots$

How to ensure that the gate does not enter hazard state?

- Environment should wait to see a change in value of output before toggling input again

Executing an ESM




Each **mode switch** corresponds to a **task**

Examples:

- Input task: $(mode = stable \ \& \ in = x) \rightarrow mode := unstable$
- Output task: $(mode = stable) \rightarrow out := x$
- Internal task: $(mode = unstable) \rightarrow \{ x := \neg x ; mode := stable \}$
- ...


Block Diagrams



Visually the same as the synchronous case

However, their execution semantics is **different!**

DoubleBuffer



$(\text{Buffer}[\text{out} \mapsto \text{temp}] \mid \text{Buffer}[\text{in} \mapsto \text{temp}]) \setminus \text{temp}$


Instantiation: Create two instances of **Buffer**

- output of **Buffer1** = input of **Buffer2** = variable **temp**

Parallel composition: Asynchronous concurrent execution of **Buffer1** and **Buffer2**

Variable hiding: Encapsulation (**temp** becomes local)

Composing Buffer1 and Buffer2





Inputs, outputs, states, and initialization for composition obtained as in synchronous case

What are the tasks of the composition?

- **Production** of output on **temp** by **Buffer1** **synchronized** with **consumption** of input on **temp** by **Buffer2**

Compiled DoubleBuffer



Asynchronous Composition (1)

Given asynchronous processes P_1 and P_2 , how to define $P_1 \mid P_2$?

In each step of execution, **only one task** is executed

- Concepts such as await-dependencies, compatibility of interfaces are not relevant

Example Case 1: (see textbook for complete definition)

If

- y is an **output** channel of P_1 and **input** channel of P_2
- A_1 is an **output** task of P_1 for y with code: $\text{Guard}_1 \rightarrow \text{Update}_1$
- A_2 is an **input** task of P_2 for y with code: $\text{Guard}_2 \rightarrow \text{Update}_2$

then

- $P_1 \mid P_2$ has an **output** task for y with code: $(\text{Guard}_1 \ \& \ \text{Guard}_2) \rightarrow \text{Update}_1 ; \text{Update}_2$

Asynchronous Composition (2)

Given asynchronous processes P_1 and P_2 , how to define $P_1 \mid P_2$?

In each step of execution, **only one task** is executed

- Concepts such as await-dependencies, compatibility of interfaces are not relevant

Example Case 2: (see textbook for complete definition)

If

- x is an **input** channel of P_1 and **input** channel of P_2
- A_1 is an **input** task of P_1 for x with code: $\text{Guard}_1 \rightarrow \text{Update}_1$
- A_2 is an **input** task of P_2 for x with code: $\text{Guard}_2 \rightarrow \text{Update}_2$

then

- $P_1 \mid P_2$ has an **input** task for x with code: $(\text{Guard}_1 \ \& \ \text{Guard}_2) \rightarrow \text{Update}_1 ; \text{Update}_2$

Asynchronous Composition (3)

Given asynchronous processes P_1 and P_2 , how to define $P_1 \mid P_2$?

In each step of execution, **only one task** is executed

- Concepts such as await-dependencies, compatibility of interfaces are not relevant

Example Case 3: (see textbook for complete definition)

If

- x is an **input** channel of P_1 and **not** a channel of P_2
- A_1 is an **input** task of P_1 for x with code: $\text{Guard}_1 \rightarrow \text{Update}_1$

then

- $P_1 \mid P_2$ has an **input** task for x with code: $\text{Guard}_1 \rightarrow \text{Update}_1$

Asynchronous Composition (4)

Given asynchronous processes P_1 and P_2 , how to define $P_1 \mid P_2$?

In each step of execution, **only one task** is executed

- Concepts such as await-dependencies, compatibility of interfaces are not relevant

Example Case 4: (see textbook for complete definition)


If

- y is an **output** channel of P_1 and **not** a channel of P_2
- A_1 is an **output** task of P_1 for y with code: $\text{Guard}_1 \rightarrow \text{Update}_1$

then

- $P_1 \mid P_2$ has an **output** task for y with code: $\text{Guard}_1 \rightarrow \text{Update}_1$

Execution Model: Another View




A single execution step:

- **Execute** an **internal task** of one of the processes, or
- **Process input** on an external channel x : execute an **input task** for x of every process to which x is an input, or
- **Execute** an **output task** for an output y of some process, followed by an **input task** for y for every process to which y is an input

If **multiple** choices are **enabled**, **choose** one **non-deterministically**

- No constraints on relative execution speeds


Asynchronous Execution




What can happen in a **single round** of this asynchronous model P ?

- $P1$ **synchronizes with the environment** to accept input on in
- $P2$ **synchronizes with the environment** to send output on out
- $P1$ **performs some internal computation** (one of its internal tasks)
- $P2$ **performs some internal computation** (one of its internal tasks)
- $P1$ **produces output** on x , followed by its immediate consumption by $P2$
- $P2$ **produces output** on y , followed by its immediate consumption by $P1$


Asynchronous Merge



$\text{Merge}[\text{out} \mapsto \text{temp}] \mid \text{Merge}[\text{in1} \mapsto \text{temp}][\text{in2} \mapsto \text{in3}]$



Asynchronous Execution



Note: Inter-process communication is *blocking*:

if no task of $P2$ associated with x is enabled in a round, $P1$ cannot write to x in that round

A process P is *non-blocking* if for every input channel in and state s of P , some task of P associated with in is enabled in state s

In designs with *non-blocking* processes, a receiving process is often expected to send an *acknowledgement* back to the sender of a message m that it did receive m

Deadlocks

A commonly occurring error in **asynchronous** designs


A **deadlock** occurs when **each process** is **waiting** for some other process to execute a task, but **no task** is **enabled**

In general, a state **s** of an asynchronous process **P** is a ***deadlock state*** if

1. **no task** is **enabled** in **s** and
2. **s** does **not** correspond to a **successful termination** of the system

The **absence** of **deadlocks** is a generic **safety requirement** that is expected to be an invariant of all asynchronous designs

Deadlocks



Process P_i can generate a request that is serviced by process P_j ($i \neq j$), which provides a response to P_i 's request

Consider this transition of $P_1 \mid P_2$ for state $(P_1.mode, y_1, P_2.mode, y_2)$:

$(idle, null, idle, null) \xrightarrow{x_1!req_1} (wait, null, idle, [req_1])$

Two tasks are enabled:

1. mode switch from $idle$ to $busy$ in P_2
2. mode switch from $idle$ to $wait$ in P_2

In the second case, we get

$(wait, null, idle, [req_1]) \xrightarrow{x_2!req_2} (wait, [req_2], idle, [req_1])$

Now, no task is enabled: **deadlock!**

Credits

Notes based on Chapter 4 of

Principles of Cyber-Physical Systems

by Rajeev Alur

MIT Press, 2015